

1 Introduction and data

In this paper we train a model to do authorship attribution. Our data is a fragment of the PAN2020 dataset ¹. The dataset contains fanfiction texts by 20 different authors. To create a model that can automatically predict whether 2 texts were written by the same author, we set up a feature extractor, followed by some classification model. By choosing our features intelligently, and processing them properly, a fairly simple classifier achieves great performance on this task.

2 Features

We represented each fanfiction snippet using 150 numerical features in seven groups. The provided training, development, and test splits contain 2,138, 268, and 267 snippets, respectively, covering 20 authors. Only the text is passed to the feature extractor: author labels are used as targets, and snippet IDs are not input features. To address contraction encoding, we normalized double quotes occurring between ASCII letters to apostrophes, then applied NLTK sentence tokenization, word tokenization and English part-of-speech (POS) tagging. Lastly, lexical counts use lowercased alphabetic tokens, while punctuation remains available in the normalized text and token sequence.

Table 1 summarizes the seven feature groups. Firstly, character and punctuation features capture surface writing conventions, while structural and lexical features describe organization and lexical preferences. Moreover, syntactic features capture grammatical patterns and narrative perspective, and informal-writing features target fanfiction conventions. At last, semantic features capture conceptual preferences and tone, but may also reflect topic and fandom.

For normalization, character counts are divided by the character count, lexical and regular-expression pattern counts by the alphabetic word count, POS counts by the total token count, and supersense counts by the number of matched noun/verb senses. Also, sentence-length measures are normalized by the number of sentences containing alphabetic words and document word

Table 1: Overview of the 150 features used to represent each fanfiction snippet.

Group	#	Description
Character	30	Letter frequencies, capitalization, digits and spacing.
Punctuation	12	Eight punctuation marks and four dialogue patterns.
Structural	8	Document size, paragraph count and sentence-length distributions.
Lexical	15	Word-length distributions, vocabulary diversity and repetition.
Syntactic	14	Eleven coarse POS frequencies and three pronoun-person frequencies.
Informal	13	Fanfiction conventions such as contractions, ellipse, stutters, etc..
Semantic	58	WordNet noun/verb supersenses, two WordNet summary measures, sentence-level VADER sentiment and ten NRC associations.
Total	150	

and sentence counts are log-transformed, whereas paragraph count is retained as a raw count. Finally, undefined ratios and empty summary statistics return zero, while StandardScaler is fitted on the training data within the classifier pipeline and scales the feature columns before classification so it does not restrict values to a fixed interval.

3 Classifier

For the classification task, we evaluated several models: a Linear Support Vector Classifier (SVC), a Logistic Regression Model (LRM), and a Random Forest (RF), all implemented using the scikit-learn library. Since many of the input features operate on different scales, and some of these models are sensitive to differences in feature magnitude, we standardized the features using scikit-learn’s StandardScaler. This transformation centers each feature around zero and scales it to unit variance. Feature scaling massively improved the F1-score on the development set, as expected.

During early tests, the models had not yet been

¹<https://pan.webis.de/clef20/pan20-web/author-identification.html>

fully tuned. To improve the performance of the final model, the next step was therefore to optimize the C parameter, which controls the degree of regularization applied to the model weights. Since both the SVC and LRM outperformed the RF model, we continued with these two models. In addition to achieving better performance, they are also simpler than the RF model. The results of this tuning process are presented in 1.




Figure 1: While the LRM is better with default c , the SVC is best with a tuned one.

4 Results

Throughout this report, we will evaluate our classifier with the macro-F1 score, the mean F1 over all authors:

$$F_1 = \frac{1}{N} \sum_{a=1}^N \frac{2P_a R_a}{P_a + R_a},$$

where P_a and R_a are the precision and recall for author a , and $N = 20$ is the total number of authors.

Standardization is essential for the linear models. Without scaling, the SVC and LRM achieved macro-F1 scores of only 0.731 and 0.386 respectively on the development set. In contrast, the RF model remained unaffected, achieving a macro-F1 score of 0.890. However, after applying the scaler and retaining the default model settings, performance improved substantially: the SVC and LRM reached macro-F1 scores of 0.923 and 0.926 respectively, outperforming the RF model. Tuning C (Figure 1) then showed that both linear models perform wbetter with much stronger regularization, as both peak at $C \approx 0.03$ (default 1). Here, the SVC reaches 0.953. We selected this model, retrained it on the training and development sets combined, and applied it once to the test set.

On the test set, the final model achieves a macro-F1 of 0.944 compared to 0.953 on the development set. This small drop from the development set is expected, since C was selected on it. Furthermore this small of a difference might be covered in noise. Specific mistakes will be discussed in section 6.




Figure 2: Confusion matrix of the final SVC on the test set, normalized per true author. Authors are labelled A–T.

5 Ablation and feature selection

The next step for us was to measure the contribution of each group. To do this, we retrain the final SVC with one group left out and evaluate it on the development set (Figure 3). We note that every group contributes. The informal group is the most informative (0.953 drops to 0.883 without it), followed by punctuation (0.909) and then character frequencies (0.915).

Feature selection would likely bring little gain. For our final model, no group is pure noise. Probably, the stronger regularization already shrinks uninformative weights. A brief test with the untuned LRM, showed that removing the semantic, structure or syntactic group still improved the score. However for our final, tuned model, this does not seem to be the case.

6 Discussion

Problems

One issue we faced was the fact that our scores were stuck being rather poor. This was resolved with the standard scalar and the non-linear RF




Figure 3: Macro-F1 of the final SVC on the development set with one feature group left out. The dark bar uses all features.

model. Another issue we encountered was the fact that the apostrophies were annotated as ; causing our initial version for those features to fail. Furthermore, the standard size of the snippets and the fact they don't use newlines means that paragraph structure features don't apply here.

6.1 Failure Analysis

On the test set, our classifier makes only 14 mistakes, for 267 snippets. To analyze these failures, we look at several different aspects. Firstly, per-author F1 ranges from 0.86 to 1.00, with no author failing badly. This means that our errors are spread out over all authors, rather than the model confusing a handful very badly, and the rest being perfect. Where the model is wrong, the margin between the score for the most likely author and second most likely author is 0.30, compared to a gap of 1.25 for correct answers. This means that the model is less confident when it makes mistakes. Finally, upon inspection of each mistakes' text, we note that 95, 141, 166, 179 all contain an author note, this might be stylistically different from the rest of that authors' text. These mistakes also had a larger margin.

With 10 extra hours

With 10 extra hours we would regroup and recreate the dataset. Currently there is leakage between train, dev and test data, as the snippets of the same stories are shared between them. This is made worse by the fact that these snippets were created with overlap, meaning that actual overlapping parts of these sequences are also shared between train and test data. We would then redo

our work on this cleaner dataset. Another thing we could tackle is trying more different features.

Challenges in authorship attribution

Authors don't always write in the same style. In our failure analysis we saw that quite some of the remaining mistakes (4/14) were made when the snippet contained some opening or closing notes, written in a different style than their normal writing. Another challenge is that the topic the authors write on might be a confounding factor, especially for some features like n-grams.

7 Conclusion

Overall, we believe that we created a classifier which, for this dataset, performs the task of authorship attribution really well. Tuning the regularization allowed us to keep all feature-groups in the final model, as the ablation study showed that each contributed at least to some degree. There were some issues with the data, which makes it hard for us to be fully confident in the exact performance. But dealing with this felt out of scope for this assignment. For future work this would be the most essential part to cover.